

Object recognition and computer vision 2020/2021

Jean Ponce, Ivan Laptev, Cordelia Schmid
and Josef Sivic

Assignment 3: Image classification

(Adapted from Rob Fergus and Soumith Chintala)

Goal

This assignment requires you to participate in a Kaggle competition with the rest of the class on the [Caltech-UCSD Birds-200-2011](#) bird dataset. The objective is to produce a model that gives the highest possible accuracy on a test dataset containing the same categories.

Guidelines

You should first clone the Github repository https://github.com/willowsierra/recvis20_a3. The *main.py* file contains code for training and evaluating your models. Once training is completed, it produces a file *kaggle.csv* that lists the IDs of the test set images, along with their predicted label. This file should be uploaded to the Kaggle webpage, which will then produce a test accuracy score. You can register for the competition using the private link: <https://www.kaggle.com/c/mva-recvis-2020/>.

What to hand-in

You should write a 1-page, double-column report in [CVPR'21 format](#) briefly presenting your approach and obtained results. This report should be uploaded in pdf format to the class Moodle by **November 24th 2020**.

The model architecture is specified in *model.py*. Currently a simple baseline model is provided. You can use this model or create your own. You are free to implement any approach covered in class, or in the research literature. Of course, tricks that you devise yourself are also encouraged. The test and training data are provided in the Kaggle competition data download page. You should unzip the images.

